

09 · Product Roadmap

A consolidated, prioritized view of **what is not yet built, partially built, or deliberately deferred** — so the Institute can plan phases, scope funding, and decide what matters. Every item is grounded in the code or the design docs (the detailed evidence and file references sit behind this summary), and each carries an honest state and a rough size.

This is a *planning* document, not a promise: sizes are estimates, and some items are questions for the Institute to decide rather than committed work.

Companion documents: [01 · Project Audit](#) (risks), [03 · Engine Comparison](#), [07 · Data Governance](#), [08 · Content Stewardship](#). Glossary in the [README](#).

0 · How to read this

State labels: **Not built** · **Partial** · **Stub** (a working seam that's switched off/faked) · **Deferred** (designed, paused on purpose) · **Known limitation** · **Config flip** (largely a setting to change).

Size: *small* (hours–days) · *medium* (weeks) · *large* (multi-week/quarter). **Type:** *engineering* vs *data/curation* vs *ops/decision*.

Two scope caveats, so the list is honest: - **Bilingual (EN/FR) and accessibility are undocumented absences, not promised work** — the app is English-only with no i18n and no stated accessibility commitment. For a Canadian public platform (with Government-of-Canada and TELUS content), we flag both as decisions to make, not deferrals already agreed. - **A few large ideas were deliberately paused as low-return** (e.g. the full grounded-corpus cutover). Those are noted as such, not as unfinished obligations.

Horizon 1 — Launch readiness (before real users and real data)

Mostly small config/ops steps, but each is a genuine gate. If the Institute does nothing else, do these.

- **Flip the web firewall (WAF) to blocking.** It's in watch-only mode today; a one-setting change after a short observation window. (*Config flip · small*)
- **Move alerts off a personal address; get SES out of sandbox.** System alerts and the weekly digest currently point at a capstone member's Northeastern email, and production email delivery is not yet enabled. Subscribe a shared Institute inbox and request SES production access. (*Ops · small*)
- **Add a custom domain** (+ certificate). None is configured today. (*Ops · small*)
- **Purge demo data and rotate the shared password** before real users — see [08 §G](#). (*Ops · small*)
- **Provision the client-owned AWS account** — its own AI (Bedrock/BDA) project + ARNs, secrets, deploy role, SES identities, and the legal-cases data. Two sandbox ARNs are hardcoded and must be replaced. See [02 · Deploy Runbook](#). (*Ops · medium*)
- **Decide the residency strategy and protect the environment.** Pick the region posture (Part Horizon 2), and enable backups/point-in-time recovery on a real Canadian environment

— today only the production stage retains data and only one table has PITR. (*Ops + decision · medium*)

- **Switch the production extraction engine to Textract-LAYOUT (+ enable OCR).** The live engine is BDA — ranked *lowest on provenance*, with inferred page numbers and no OCR for scanned PDFs — while every handoff doc recommends Textract-LAYOUT. Largely a config-and-revalidate change that directly improves the citations users see. See 03. (*Config flip + validation · small-medium*)
-

Horizon 2 — The trust backbone (make identity and “confirmed” real)

The platform’s core value — a *credible* index — rests on verified identity and a working confirmation path. Several pieces are stubbed or structurally incomplete.

- **Activate business-registry (ISED) verification.** The seam exists but runs as a **stub**; until it’s live, “registry-verified” means “matched a built-in list.” Confirm the live API and switch it on. This unblocks self-serve upload and real supplier verification. (*Engineering · small-medium*)
 - **Finish the Business-Number crosswalk curation** (37 of 103 orgs done; the file says “*verify before the prod migration*”). Human registry lookups; gated on the ISED check above for auto-validation. (*Data/curation · medium*)
 - **Fix or honestly reframe the “Confirmed” tier.** On the Index it is **structurally always 0%** — no runtime path can move a commitment to “confirmed.” Either build the counterparty-attestation bridge or reframe the tile so it isn’t silently empty. (*Engineering + data · medium*)
 - **Onboard real Indigenous suppliers.** The supplier confirmation *flow* is real code, but the supplier pool and confirmations are fixtures (“until a real set of suppliers is seeded, matches are fixture-quality”). This is an adoption/data effort more than a build. (*Data + adoption · medium-large*)
 - **Build the company-account ↔ seeded-data crosswalk** (flagged internally as **P0**). Today a logged-in company is linked to a separate data world from its own public seeded commitments — the two don’t join. (*Engineering · medium*)
 - **Ship the “trackability” signal** (a rare quick win). The data model already records whether each commitment has a due date and a measurable target; surfacing *how many of an org’s commitments are actually trackable* turns the honest gap in 06 into a first-class accountability feature. (*Engineering · small*)
 - **Build the deferred governance enforcement layer** — the consent record, access-audit log, **ca-central-1** migration, customer-managed encryption key, and per-data-class access controls. Designed and approved; **only the classification tag is built** so far. Required before ingesting any private or community-linked data. See 07. (*Engineering + governance · large*)
-

Horizon 3 — Reach, depth, and the bigger product

- **Bilingual EN/FR** — no internationalization exists; the UI is English-only. An official-languages question for a Canadian public platform. (*Engineering + translation · large · decision needed*)

- **Accessibility (WCAG) audit and remediation** — no stated commitment today. (*Engineering + audit · medium · decision needed*)
- **RAP maturity / completeness scoring + feedback report** — the largest unbuilt *product* idea: score a RAP’s depth against an Indigenomics standard and emit a structured feedback report. Needs a rubric co-defined with the Institute. (*Engineering + data · large*)
- **Real authentication** (replace the shared demo password with Cognito-style logins + password reset). (*Engineering · medium*)
- **Real supplier identity verification** (CCAB / nation endorsement) to replace the stubbed identity tier — today a tier is just an enum. (*Engineering + partnership · large*)
- **Self-serve org upload for unclaimed organizations** (today only claimed companies can upload; gated on registry verification being real). (*Engineering · medium*)
- **Alignment engine depth** — semantic matching and AI rationale are **stubbed off by default** (BM25 + templates run today); plus equity/capital matching, a track-record signal, a feedback loop, and a proactive inbox. (*Engineering · medium-large*)
- **Web / statutory RAP ingestion** (HTML and statutory sources, plus a discovery crawler — RAP discovery is manual today). (*Engineering · medium*)
- **Legal-info depth** — multi-turn Q&A and a verified clinic directory; retrieval enhancements (rerank/ANN/snippets) are deliberately trigger-gated until scale warrants them. (*Engineering + content · medium*)
- **Whole-product Horizon-2 gaps** — dispute-resolution workflow, an Institute admin/curation surface, a visibility/privacy model, annual re-confirmation lifecycle. (*Engineering · large, multi-quarter*)

Quick reference — the backlog by theme

#	Item	State	Size / type
Data & Trust	ISED registry verification	Stub	small-med · eng
	BN-map curation (37/103)	Partial	med · data
	“Confirmed” tier (always 0%)	Not built	med · eng+data
	Real supplier onboarding	Partial (code real, data fixture)	med-lg · data
	Trackability signal	Not built	small · eng (quick win)
	Governance enforcement (consent/audit/KMS/central-1)	Deferred (tag only built)	large · eng+gov
Extraction	Prod engine → Textract-LAYOUT	Config flip	small-med · eng
	Scanned-PDF OCR	Known limitation	small-med · eng
	In-country (“Option B”) extraction	Deferred	med · eng+decision

#	Item	State	Size / type
Security & Ops	TELUS Canada-hosted inference	Prospective	med · eng+partner
	RAP maturity scoring + feedback	Not built	large · eng+data
	At-scale run + job queue	Known limitation	med · eng
	WAF → blocking	Config flip	small
	Custom domain	Known limitation	small
	Shared-inbox alerts + SES prod access	Partial	small · ops
	Client-account provisioning	To do	med · ops
	Backups/PITR + residency decision	Partial	med · ops+decision
	Infra cost wins (arm64, lifecycle, TTL)	Not done	small · config
	Bilingual EN/FR	Absent	large · eng+content
Reach & Access	Accessibility (WCAG)	Absent	med · eng+audit
	Self-serve org upload	Partial	med · eng
	Real auth (Cognito)	Deferred	med · eng
	Supplier identity verification	Stub	large · eng+partner
	Alignment engine depth	Stub-by-default + planned	med–lg · eng
Product	Company ↔ seeded-data crosswalk	Not built (P0)	med · eng
	Dispute/admin/visibility (Horizon 2)	Not built	large · eng
	Legal Q&A + retrieval depth	Deferred	med · eng+content

Deliberately paused (not recommended now): the full **grounded-corpus cutover** was evaluated and paused as low-return — the live Index stays on the curated illustrative dataset. The **second-AI “judge” validation** is intentionally omitted, because the research showed it gives no reliable signal on this data (06); human review stays the right instrument.

Suggested first five (if you must pick)

1. **Launch-readiness bundle** (Horizon 1): WAF→block, custom domain, shared-inbox alerts + SES, demo purge. Small steps, but real gates — and the alert email currently points at an address that will be deactivated.

2. **Switch production to Textract-LAYOUT + OCR.** Largely config; directly improves the citations every user sees.
3. **Activate ISED registry + finish BN curation.** The trust backbone; both are self-flagged blockers with the seam already built.
4. **The trackability signal.** Highest visibility per unit of effort — the data already exists.
5. **Company ↔ seeded-data crosswalk + reframe the 0% “Confirmed” tier.** Two quiet value-proposition holes worth closing (or honestly reframing) early.

*Grounded in the repository as handed off: the design specs’ Future/Out-of-scope sections, the internal action-items evaluation, and the code seams (`registry.ts`, `org-bn-map.ts`, the WAF and engine config in `sst.config.ts`, `insights.ts`). Where an item is a **decision** rather than a build — residency posture, bilingual support, accessibility target, the RAP-maturity rubric — it is flagged as such, because those are the Institute’s to make.*