

11 · Licensing & Intellectual Property

When a client takes ownership of a platform, they need to know **what they own, what they're licensed to use, and what obligations come attached**. Today the repository does **not** answer that cleanly — there is no license file and no license declared — so this document sets out the current state and the decisions to make before the platform is distributed or launched.

This is an informational summary, **not legal advice**. The Institute should have counsel confirm the license choice and any data/trademark questions before public launch.

Companion documents: [01 · Project Audit](#), [02 · Deploy Runbook](#), [07 · Data Governance](#). Glossary in the [README](#).

Part A · The code — ownership and license (decision needed)

Current state: the repository has **no LICENSE file**, and `package.json` sets `"private": true` with **no license field**. In practice that means the code is “**all rights reserved**” by default and not licensed to anyone — which is fine for a private handoff to one client, but it leaves ownership and reuse **undefined in writing**.

Decisions to make (with the team and, if relevant, the university):

- **Who owns the copyright?** Capstone/student work, university, the Institute, or shared? This should be stated explicitly in a short IP-assignment or handoff letter — verbal understanding is not enough for an asset the Institute will operate publicly.
- **What license governs the client's use and any future distribution?** Options, roughly:
 - *Private / all-rights-reserved to the Institute* — simplest; the Institute owns and uses it, no public reuse. Keep `private: true`, add a short proprietary LICENSE naming the owner.
 - *A permissive open-source license* (MIT / Apache-2.0) — if the intent is to let others reuse or contribute (common for mission-driven civic tech). Requires the copyright holders' agreement.
 - *A source-available / non-commercial license* — if the Institute wants visibility but not commercial reuse by others.
- **Add the chosen LICENSE file** (and, if open-sourcing, a NOTICE/attribution file). This is a small task once the decision is made.

Recommendation: decide ownership first (an IP-assignment/handoff letter), then add a matching LICENSE file before any public repository or distribution.

Part B · Third-party dependencies (low risk, but audit before distribution)

The platform's **21 direct runtime dependencies** are the usual, permissively-licensed building blocks:

- **AWS SDK** (@aws-sdk/*, aws-xray-sdk-core) — Apache-2.0
- **Next.js, React, React-DOM** — MIT
- **Nivo charts** (@nivo/*) — MIT

- **pdf-lib, pdf-parse, robots-parser** — MIT

These are all **permissive** (MIT / Apache-2.0) — no copyleft (GPL-style) obligations that would force the Institute to open-source its own code, and no per-seat commercial licensing. Apache-2.0 carries a lightweight attribution/NOTICE obligation if the code is redistributed.

One caveat: this is a review of *direct* dependencies only. The full **transitive** dependency tree (hundreds of packages) has not been machine-audited. Before distributing the code or publishing the repository, run a license scan — e.g. `npx license-checker --summary` — to confirm nothing copyleft or restrictive has slipped in transitively, and to generate an attributions list. No license-checker tool is currently in the repo; adding one to the pre-distribution checklist is prudent.

Part C · Data & content IP

The platform involves several distinct bodies of content, each with its own IP posture:

- **The RAP Index data (curated commitments).** These are **facts drawn from organizations' public disclosures**, each stored with a link back to its first-party source. Factual data compiled from public sources is generally low-risk to index, but **how the Institute republishes it** (verbatim excerpts vs. summarized facts, with source attribution) is worth a deliberate policy — the platform already stores and displays the source link for every entry (see [10 · Methodology](#)).
 - **Legal cases corpus.** Public **court decisions** — public records by construction. Any per-source reproduction terms (some court sites impose crawling/reuse conditions) should be respected; the platform already honors `robots.txt` (`robots-parser` dependency).
 - **Sample / test fixtures — must be removed before distribution.** The repository's test data (`scripts/fixtures/`) contains derived JSON, some of which holds **verbatim text excerpts from real RAPs** (RBC and Bank of Canada page excerpts in the `textlayer-geometry-*` fixtures) used only for development. These must be purged before real deployment or any public repository — this is already on the pre-launch checklist ([02 · Deploy Runbook §5.3](#), [08 · Content Stewardship §G](#)).
 - **Institute-generated content** (framing, methodology, briefings) — owned by the Institute.
-

Part D · AI model & cloud service terms

The platform's AI features run on **Amazon Bedrock**, invoking third-party and Amazon models (Anthropic Claude, Meta Llama, Amazon Titan/BDA). Two things to be aware of:

- **Usage is governed by AWS's and the model providers' terms.** When the Institute runs the platform in its **own** AWS account, it accepts those terms directly. Model availability, pricing, and acceptable-use are set by AWS/the providers, not by this code.
 - **Output ownership.** Under current AWS/Anthropic terms, customers generally own the inputs and outputs of their model calls, but the Institute should confirm the current terms for its account — particularly if outputs (extracted commitments, briefings) are republished.
 - See [07 · Data Governance](#) for the residency implications of where that inference runs.
-

Part E · Trademarks

- **OCAP® is a registered trademark of the First Nations Information Governance Centre (FNIGC).** The platform uses the mark (with the ® symbol) to describe alignment with the OCAP principles. Continue to use it correctly and attributed; if the platform ever *claims* OCAP compliance or certification (as opposed to describing alignment with the principles), confirm the appropriate usage/permission with FNIGC. Note that 07 is candid that the OCAP *enforcement* layer is partly designed-not-built — so describe alignment, not certification.
 - **Third-party names** in the data (company names, “CCAB”, “PAR”, organization marks) are used nominatively to identify real organizations and their programs — standard for a public index, but keep usage factual and non-endorsing (see 10 · [Methodology](#)).
 - **Institute branding** (name, logo) is the Institute’s own.
-

Part F · Pre-distribution IP checklist

Before the platform is published, distributed, or launched publicly:

- ☐ **Decide and document code ownership** (IP-assignment / handoff letter).
 - ☐ **Add a LICENSE file** matching that decision (+ NOTICE if open-sourcing).
 - ☐ **Run a dependency license scan** (`npm license-checker`) and keep the attribution output.
 - ☐ **Remove the sample fixtures containing verbatim RAP excerpts** from `scripts/fixtures/` (also a data-hygiene step in 08 §G).
 - ☐ **Confirm AWS/Bedrock model terms** for the Institute’s own account.
 - ☐ **Confirm OCAP® usage posture** (alignment, not certification) and any court-source reuse terms.
 - ☐ **Have counsel review** the above and the [methodology/right-of-reply](#) posture.
-

This document reflects the repository as handed off: no license is declared today (`package.json` “private”: true, no LICENSE file), direct dependencies are permissively licensed, and the only known copyright-sensitive content — real RAP PDFs used as test fixtures — is already flagged for removal before launch. The main action is a deliberate ownership-and-license decision, which only the team and Institute can make.